Limitations of Transfer Under Distribution Drift in QA

Abstract

We investigate failure modes of zero-shot question answering using RoBERTa-large on SQuAD-v2. We report significant performance degradation under raw corpus evaluation.

Methodology and Experiments

4. Experimental Protocol
We evaluate across the comprehensive uncurated corpus without post-hoc length filtering. We configure learning rate lr=3e-5 with batch size 32 using adamw optimizer for 3 epochs.

5. Empirical Results
Under strict evaluation conditions, the framework demonstrates brittle generalization reaching Macro-F1 of 76.8% ± 0.5%.